

Checking a Casino



Fair coin:
 $\text{Pr}(\text{Heads}) = 0.5$



Biased coin:
 $\text{Pr}(\text{Heads}) = 0.75$



Suppose either a fair or biased coin was used to generate a sequence of heads & tails. But we don't know which type of coin was actual used.

Heads/Tails: ↑ ↑ ↓ ↓ ↓ ↓ ↑ ↑ ↑ ↑ ↓ ↑ ↓ ↑ ↓ ↑

Checking a Casino



Fair coin:
 $\text{Pr}(\text{Heads}) = 0.5$



Biased coin:
 $\text{Pr}(\text{Heads}) = 0.75$



Suppose either a fair or biased coin was used to generate a sequence of heads & tails. But we don't know which type of coin was actual used.

Heads/Tails: ↑ ↑ ↓ ↓ ↓ ↓ ↑ ↑ ↑ ↑ ↓ ↑ ↓ ↑ ↓ ↑

Checking a Casino



Fair coin:
 $\text{Pr}(\text{Heads}) = 0.5$



Biased coin:
 $\text{Pr}(\text{Heads}) = 0.75$



Suppose either a fair or biased coin was used to generate a sequence of heads & tails. But we don't know which type of coin was actual used.

Heads/Tails: ↑ ↑ ↓ ↓ ↓ ↓ ↑ ↑ ↑ ↑ ↓ ↑ ↓ ↑ ↓ ↑

How could we guess which coin was more likely?

Compute the Probability of the Observed Sequence

Fair coin: $\Pr(\text{Heads}) = 0.5$
Biased coin: $\Pr(\text{Heads}) = 0.75$

$x =$       

$\Pr(x \mid \text{Fair}) =$ 0.5 0.5 0.5 0.5 0.5 0.5 0.5

$\Pr(x \mid \text{Biased}) =$ 0.75 0.75 0.25 0.25 0.25 0.25 0.75

Compute the Probability of the Observed Sequence

Fair coin: $\Pr(\text{Heads}) = 0.5$
Biased coin: $\Pr(\text{Heads}) = 0.75$

$x =$       

$$\Pr(x \mid \text{Fair}) = 0.5 \times 0.5 \times 0.5 \times 0.5 \times 0.5 \times 0.5 \times 0.5 = 0.5^7 = 0.0078125$$

$$\Pr(x \mid \text{Biased}) = 0.75 \times 0.75 \times 0.25 \times 0.25 \times 0.25 \times 0.25 \times 0.75 = 0.001647949$$

Compute the Probability of the Observed Sequence

Fair coin: $\Pr(\text{Heads}) = 0.5$
Biased coin: $\Pr(\text{Heads}) = 0.75$

$x =$ 

$$\Pr(x \mid \text{Fair}) = 0.5 \times 0.5 \times 0.5 \times 0.5 \times 0.5 \times 0.5 \times 0.5 = 0.5^7 = 0.0078125$$

$$\Pr(x \mid \text{Biased}) = 0.75 \times 0.75 \times 0.25 \times 0.25 \times 0.25 \times 0.25 \times 0.75 = 0.001647949$$


The *log-odds* score:

$$\log_2 \frac{\Pr(x \mid \text{Fair})}{\Pr(x \mid \text{Biased})} = \log_2 \frac{0.0078}{0.0016} = 2.245$$

> 0 . Hence “Fair” is a better guess.

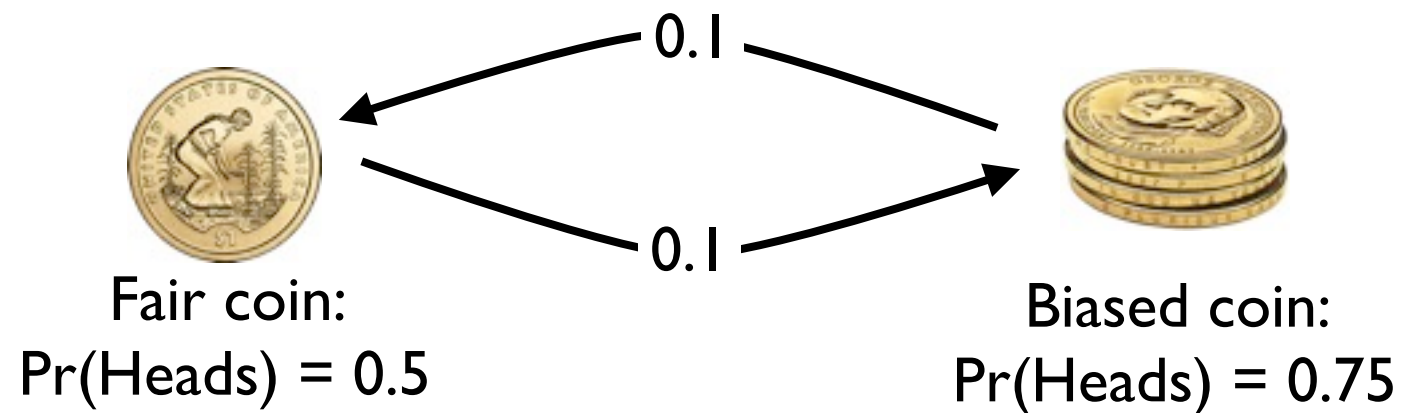
What if the casino switches coins?

Fair coin: $\Pr(\text{Heads}) = 0.5$
Biased coin: $\Pr(\text{Heads}) = 0.75$
Probability of switching coins = 0.1



What if the casino switches coins?

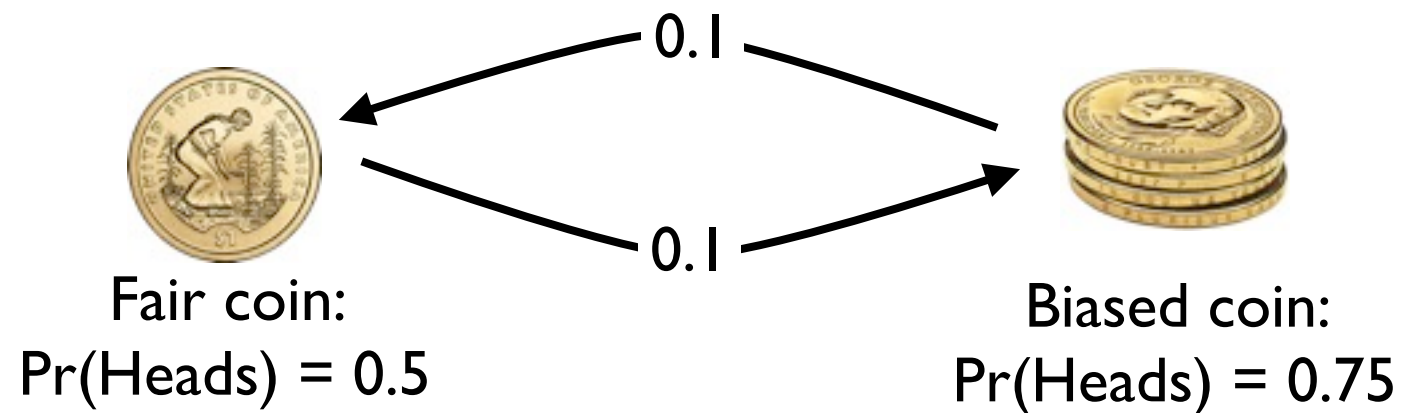
Fair coin: $\Pr(\text{Heads}) = 0.5$
Biased coin: $\Pr(\text{Heads}) = 0.75$
Probability of switching coins = 0.1



How can we compute the probability of the entire sequence?

What if the casino switches coins?

Fair coin: $\Pr(\text{Heads}) = 0.5$
Biased coin: $\Pr(\text{Heads}) = 0.75$
Probability of switching coins = 0.1



How can we compute the probability of the entire sequence?

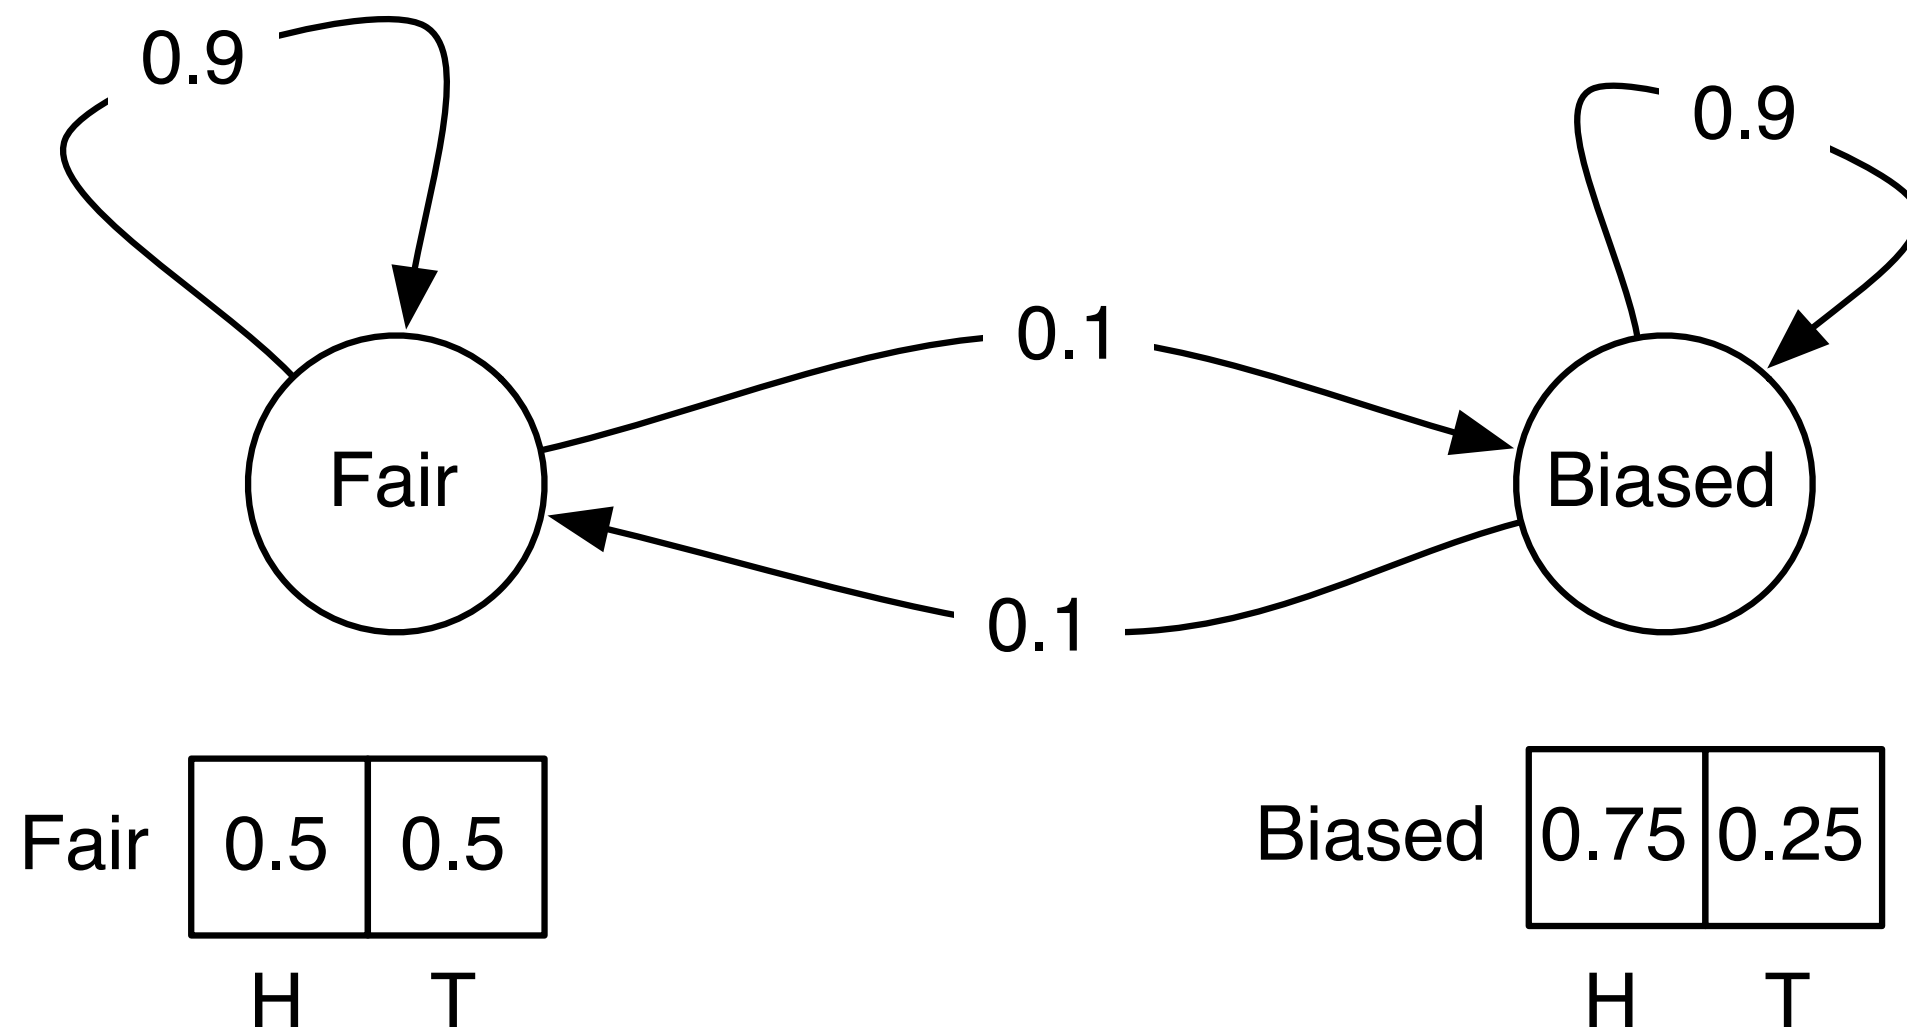
How could we guess which coin was more likely **at each position**?

Hidden Markov Model (HMM)

Fair coin: $\Pr(\text{Heads}) = 0.5$

Biased coin: $\Pr(\text{Heads}) = 0.75$

Probability of switching coins = 0.1



Formal Definition of a HMM

Σ = alphabet of symbols.

Q = set of states.

A = an $|Q| \times |Q|$ matrix where entry (k,l) is the probability of moving from state k to state l .

E = a $|Q| \times |\Sigma|$ matrix, where entry (k,b) is the probability of emitting b when in state k .

$A =$

7							
6							
5							
4							
3							
2							
1							
	1	2	3	4	5	6	7



Probability of going from state 5 to state 7

$E =$

7				
6				
5				
4				
3				
2				
1				
	A	C	G	T

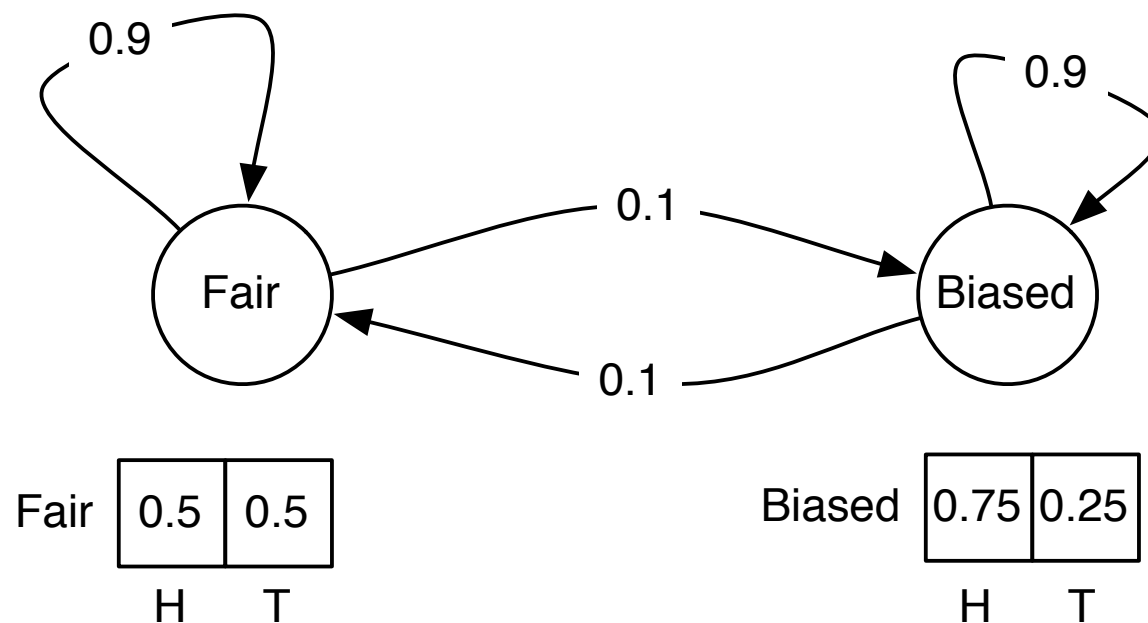
Probability of emitting T when in state 4.

Constraints on A and E



Sum of the # in each row must be 1.

Computing Probabilities Given Path



$x =$ ↓ ↑ ↓ ↑ ↑ ↑ ↓ ↑ ↑ ↓

$\pi =$ F F F B B B B F F F

$\Pr(x_i \mid \pi_i) =$ 0.5 0.5 0.5 0.75 0.75 0.75 0.25 0.5 0.5 0.5

$\Pr(\pi_i \rightarrow \pi_{i+1}) =$ 0.1 0.9 0.9 0.1 0.9 0.9 0.9 0.1 0.1 0.1

The Decoding Problem

Given x and π , we can compute:

- $\Pr(x \mid \pi)$: product of $\Pr(x_i \mid \pi_i)$
- $\Pr(\pi)$: product of $\Pr(\pi_i \rightarrow \pi_{i+1})$
- $\Pr(x, \pi)$: product of all the $\Pr(x_i \mid \pi_i)$ and $\Pr(\pi_i \rightarrow \pi_{i+1})$

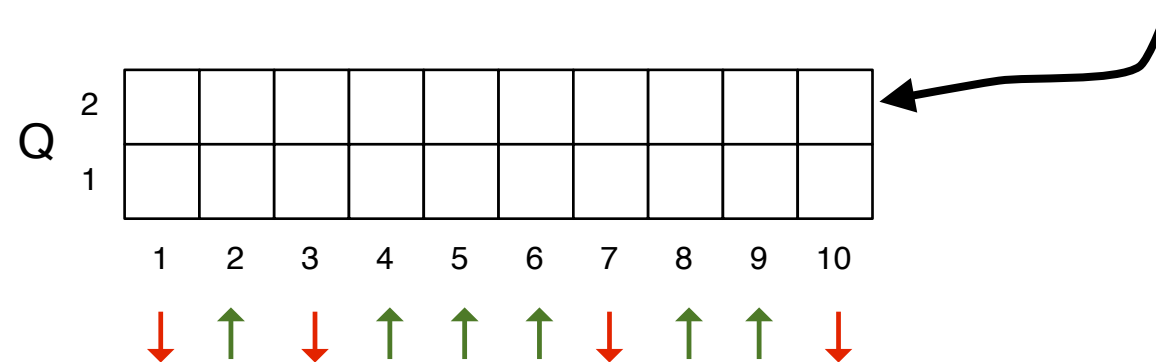
$$\Pr(x, \pi) = \Pr(\pi_0 \rightarrow \pi_1) \prod_{i=1}^n \Pr(x_i \mid \pi_i) \Pr(\pi_i \rightarrow \pi_{i+1})$$

But they are “hidden” Markov models because π is unknown.

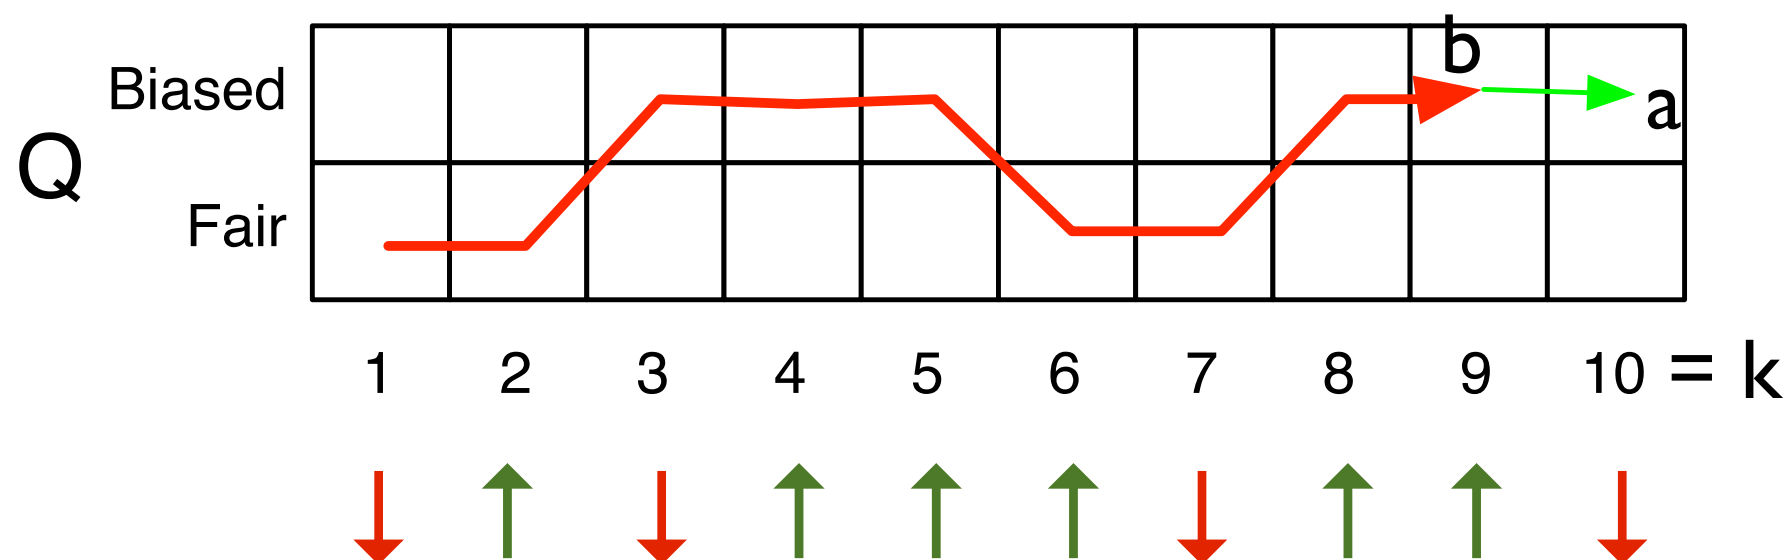
Decoding Problem: Given a sequence $x_1x_2x_3\dots x_n$ generated by an HMM (Σ, Q, A, E) , find a path π that maximizes $\Pr(x, \pi)$.

The Viterbi Algorithm to Find Best Path


$A[a, k] :=$ the probability of the **best** path for $x_1 \dots x_k$ that ends at state a .



$A[a, k] =$ the path for $x_1 \dots x_{k-1}$ that goes to some state b times cost of a transition from b to i , and then to output x_k from state a .



Viterbi DP Recurrence



$$A[a, k] = \max_{b \in Q} \{ \underbrace{A[b, k-1]}_{\text{Best path for } x_1 \dots x_k \text{ ending in state } b} \times \underbrace{\text{Pr}(b \rightarrow a)}_{\text{Probability of transitioning from state } b \text{ to state } a} \times \underbrace{\text{Pr}(x_k \mid \pi_k = a)}_{\text{Probability of outputting } x_k \text{ given that the } k\text{th state is } a} \}$$

Over all possible previous states.

Best path for $x_1 \dots x_k$ ending in state b

Probability of transitioning from state b to state a

Probability of outputting x_k given that the k th state is a .

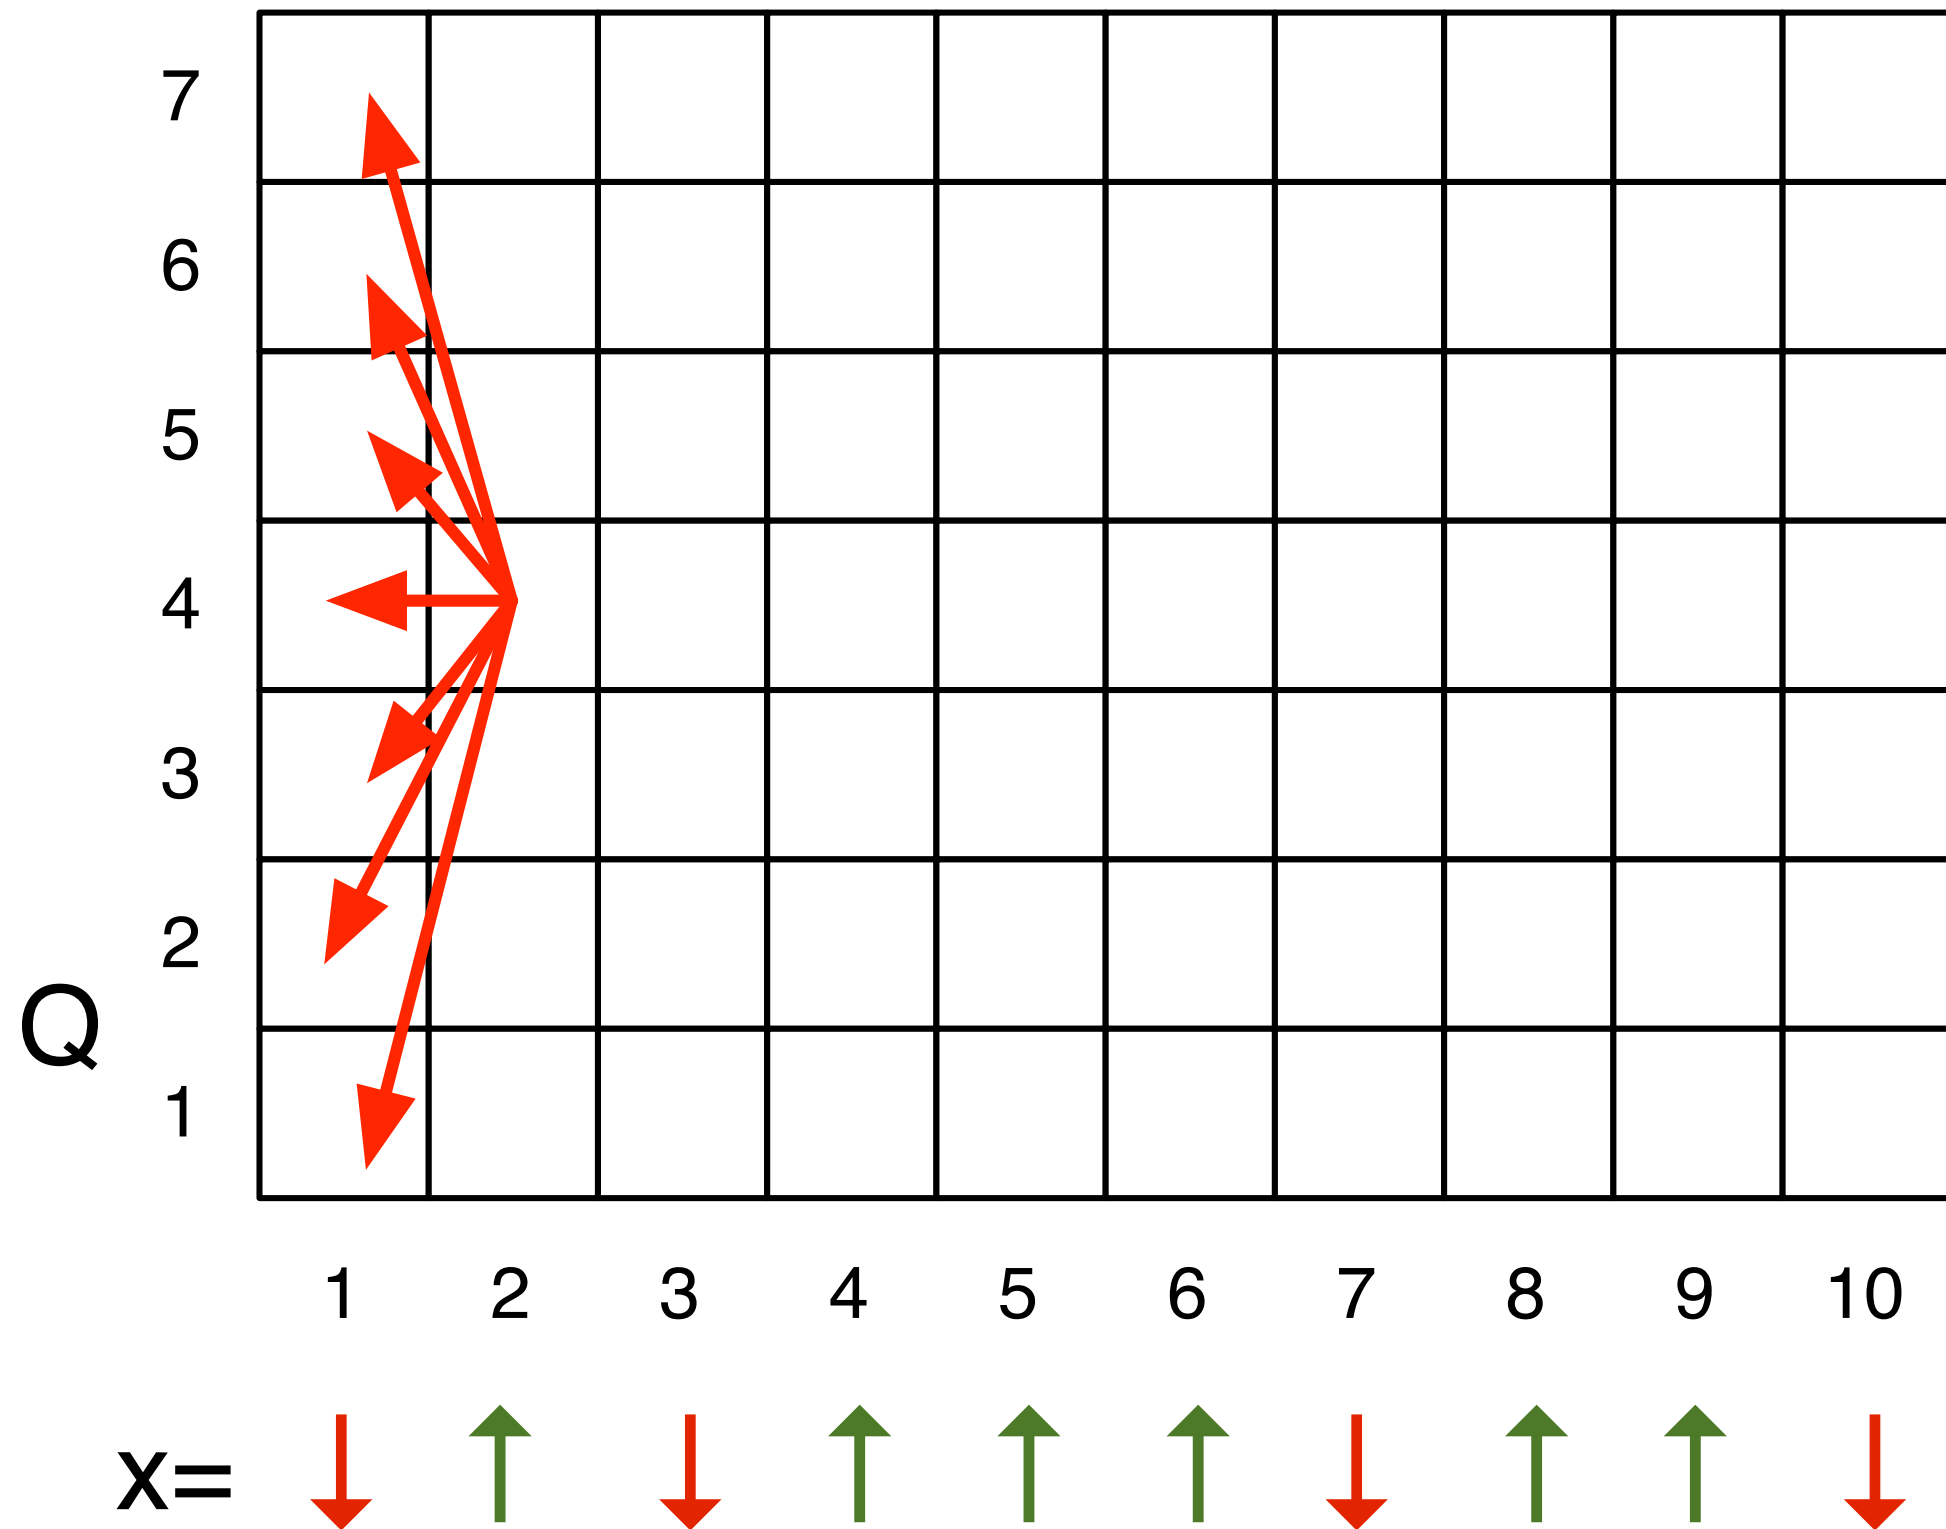
Base case:

$$A[a, 1] = \underbrace{\text{Pr}(\pi_1 = a)}_{\text{Probability that the first state is } a} \times \underbrace{\text{Pr}(x_1 \mid \pi_1 = a)}_{\text{Probability of emitting } x_1 \text{ given the first state is } a}$$

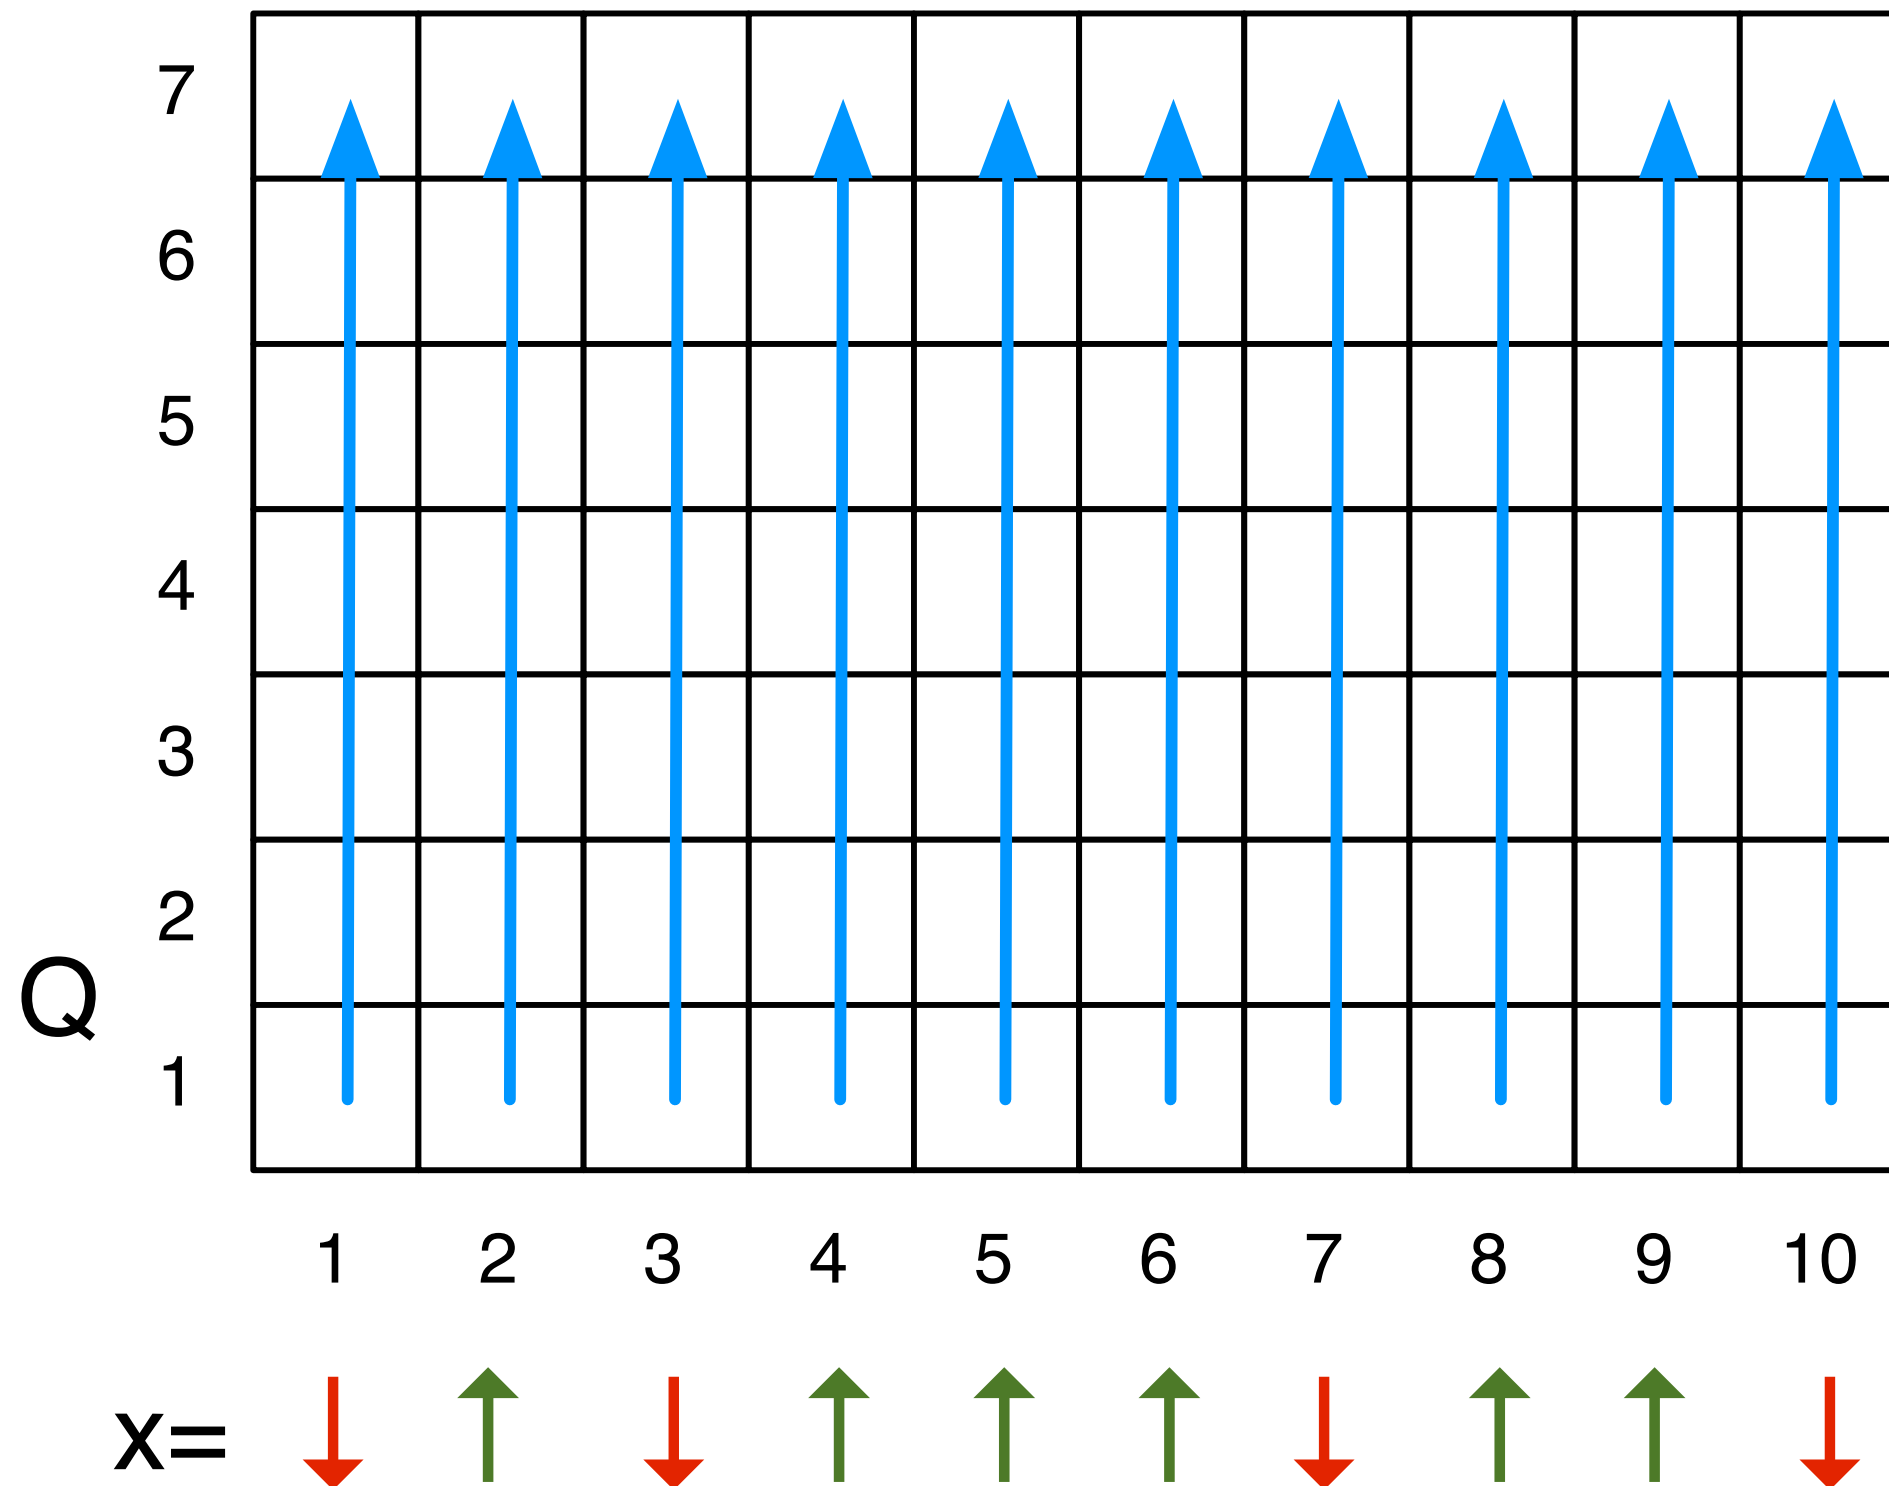
Probability that the first state is a

Probability of emitting x_1 given the first state is a .


Which Cells Do We Depend On?




Order to Fill in the Matrix:



Where's the answer?



Graph View of Viterbi



Running Time

- # of subproblems = $O(n|Q|)$, where n is the length of the sequence.
- Time to solve a subproblem = $O(|Q|)$
- Total running time: $O(n|Q|^2)$

Using Logs

Typically, we take the log of the probabilities to avoid multiplying a lot of terms:

$$\begin{aligned}\log(A[a, k]) &= \max_{b \in Q} \{\log(A[b, k-1] \times \Pr(b \rightarrow a) \times \Pr(x_k \mid \pi_k = a))\} \\ &= \max_{b \in Q} \{\log(A[b, k-1]) + \log(\Pr(b \rightarrow a)) + \log(\Pr(x_k \mid \pi_k = a))\}\end{aligned}$$

Remember: $\log(ab) = \log(a) + \log(b)$
--

Why do we want to avoid multiplying lots of terms?

Using Logs

Typically, we take the log of the probabilities to avoid multiplying a lot of terms:

$$\begin{aligned}\log(A[a, k]) &= \max_{b \in Q} \{\log(A[b, k-1] \times \Pr(b \rightarrow a) \times \Pr(x_k \mid \pi_k = a))\} \\ &= \max_{b \in Q} \{\log(A[b, k-1]) + \log(\Pr(b \rightarrow a)) + \log(\Pr(x_k \mid \pi_k = a))\}\end{aligned}$$

Remember: $\log(ab) = \log(a) + \log(b)$

Why do we want to avoid multiplying lots of terms?

Multiplying leads to very small numbers:

$$0.1 \times 0.1 \times 0.1 \times 0.1 \times 0.1 = 0.00001$$

This can lead to underflow.

Taking logs and adding keeps numbers bigger.

Estimating HMM Parameters

$$\begin{aligned}
 (\mathbf{x}^{(1)}, \boldsymbol{\pi}^{(1)}) &= \left. \begin{array}{cccccc} x_1^{(1)} & x_2^{(1)} & x_3^{(1)} & x_4^{(1)} & x_5^{(1)} & \dots & x_n^{(1)} \\ \pi_1^{(1)} & \pi_2^{(1)} & \pi_3^{(1)} & \pi_4^{(1)} & \pi_5^{(1)} & \dots & \pi_n^{(1)} \end{array} \right\} \\
 (\mathbf{x}^{(2)}, \boldsymbol{\pi}^{(2)}) &= \left. \begin{array}{cccccc} x_1^{(2)} & x_2^{(2)} & x_3^{(2)} & x_4^{(2)} & x_5^{(2)} & \dots & x_n^{(2)} \\ \pi_1^{(2)} & \pi_2^{(2)} & \pi_3^{(2)} & \pi_4^{(2)} & \pi_5^{(2)} & \dots & \pi_n^{(2)} \end{array} \right\}
 \end{aligned}$$

Training examples where outputs and paths are known.

of times transition
 $a \rightarrow b$ is observed.

$$\Pr(a \rightarrow b) = \frac{A_{ab}}{\sum_{q \in Q} A_{aq}}$$

of times x was
 observed to be
 output from state a .

$$\Pr(x \mid a) = \frac{E_{xa}}{\sum_{x \in \Sigma} E_{xq}}$$

Pseudocounts

of times transition
 $a \rightarrow b$ is observed.

$$\Pr(a \rightarrow b) = \frac{A_{ab}}{\sum_{q \in Q} A_{aq}}$$

of times x was
observed to be
output from state a .

$$\Pr(x \mid a) = \frac{E_{xa}}{\sum_{x \in \Sigma} E_{xa}}$$

What if a transition or emission is never observed in the training data?
 \Rightarrow 0 probability

Meaning that if we observe an example with that transition or emission in the real world, we will give it 0 probability.

But it's unlikely that our training set will be large enough to observe every possible transition.

Hence: we take $A_{ab} = (\text{\#times } a \rightarrow b \text{ was observed}) + 1$ \leftarrow “pseudocount”
Similarly for E_{xa} .

Viterbi Training

- **Problem:** typically, in the real world we only have examples of the output x , and we don't know the paths π .

Viterbi Training Algorithm:

1. Choose a random set of parameters.
2. Repeat:
 1. Find the best paths.
 2. Use those paths to estimate new parameters.

This is an local search algorithm.

It's also an example of a “Gibbs sampling” style algorithm.

The Baum-Welch algorithm is similar, but doesn't commit to a single best path for each example.

Some probabilities we are interested in

What is the probability of observing a string x under the assumed HMM?

$$\Pr(x) = \sum_{\pi} \Pr(x, \pi)$$

What is the probability of observing x using a path where the i^{th} state is a ?

$$\Pr(x, \pi_i = a) = \sum_{\pi: \pi_i = a} \Pr(x, \pi)$$

What is the probability that the i^{th} state is a ?

$$\Pr(\pi_i = a | x) = \frac{\Pr(x, \pi_i = a)}{\Pr(x)}$$